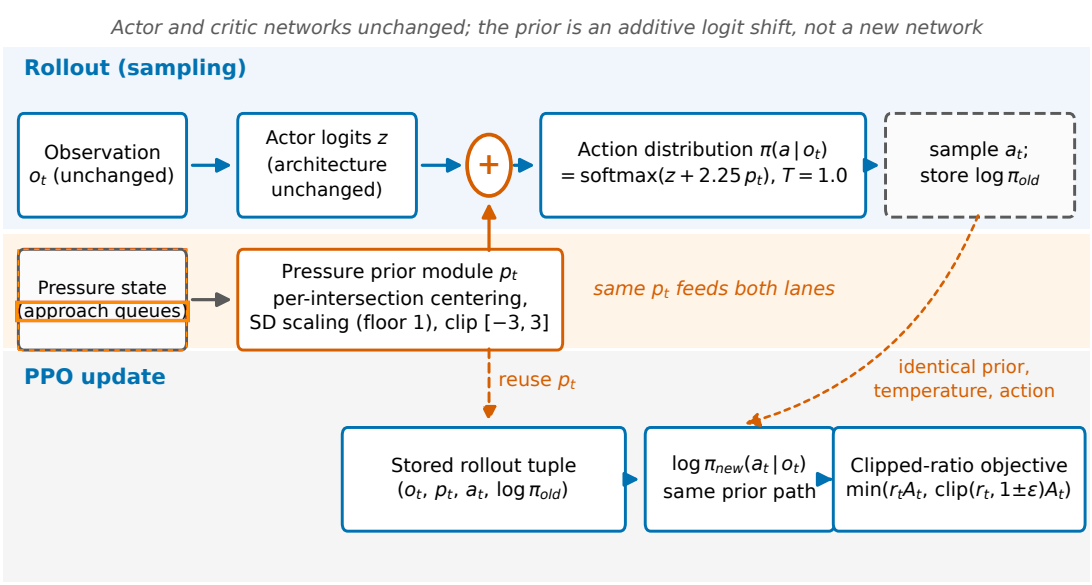
Actor and critic networks unchanged; the prior is an additive logit shift, not a new network
Rollout (sampling)
Observation o_t (unchanged)
Actor logits z (architecture unchanged)
+
Action distribution $\pi(a \mid o_t)$ = softmax$(z + 2.25\, p_t)$, $T = 1.0$
sample a_t; store $\log \pi_{old}$
Pressure state (approach queues)
Pressure prior module p_t per-intersection centering, SD scaling (floor 1), clip $[-3, 3]$
same p_t feeds both lanes
PPO update
reuse p_t
identical prior, temperature, action
Stored rollout tuple $(o_t,\ p_t,\ a_t,\ \log \pi_{old})$
$\log \pi_{new}(a_t \mid o_t)$ same prior path
Clipped-ratio objective $\min(r_t A_t,\ \text{clip}(r_t, 1 \pm \varepsilon) A_t)$