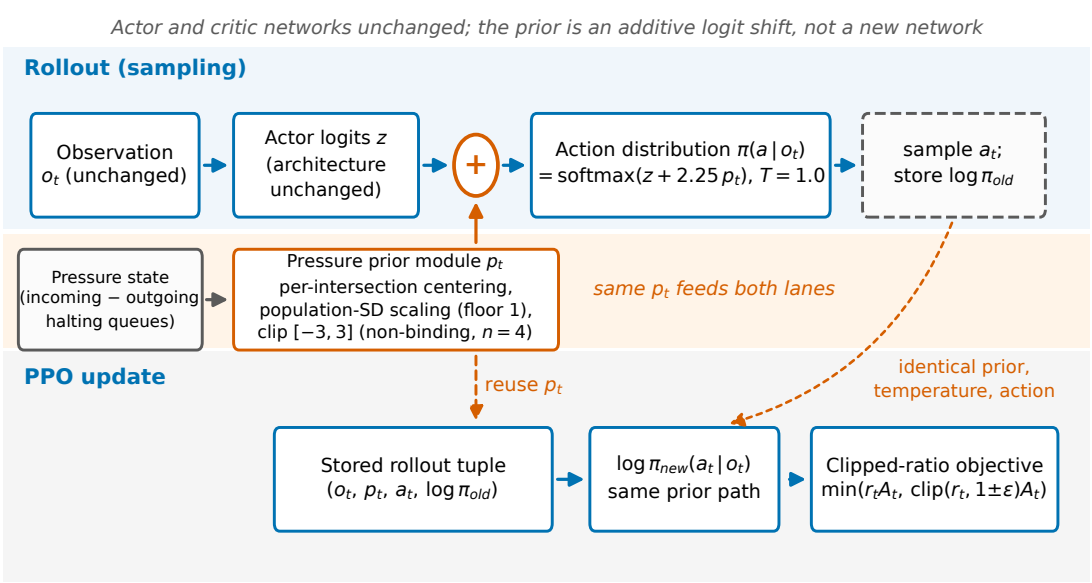
Actor and critic networks unchanged; the prior is an additive logit shift, not a new network
Rollout (sampling)
Observation o_t (unchanged)
Actor logits z (architecture unchanged)
+
Action distribution $\pi(a \mid o_t)$ = softmax($z + 2.25\,p_t$), $T = 1.0$
sample a_t; store $\log \pi_{old}$
Pressure state (incoming − outgoing halting queues)
Pressure prior module p_t per-intersection centering, population-SD scaling (floor 1), clip $[-3, 3]$ (non-binding, $n = 4$)
same p_t feeds both lanes
PPO update
reuse p_t
identical prior, temperature, action
Stored rollout tuple (o_t, p_t, a_t, $\log \pi_{old}$)
$\log \pi_{new}(a_t \mid o_t)$ same prior path
Clipped-ratio objective $\min(r_t A_t,\ \mathrm{clip}(r_t, 1 \pm \varepsilon) A_t)$